Quality–latency trade-off by bucket count

Qwen3.8-27B · CUDA Graph · 128K input · 8,193 decode steps · batch size 1

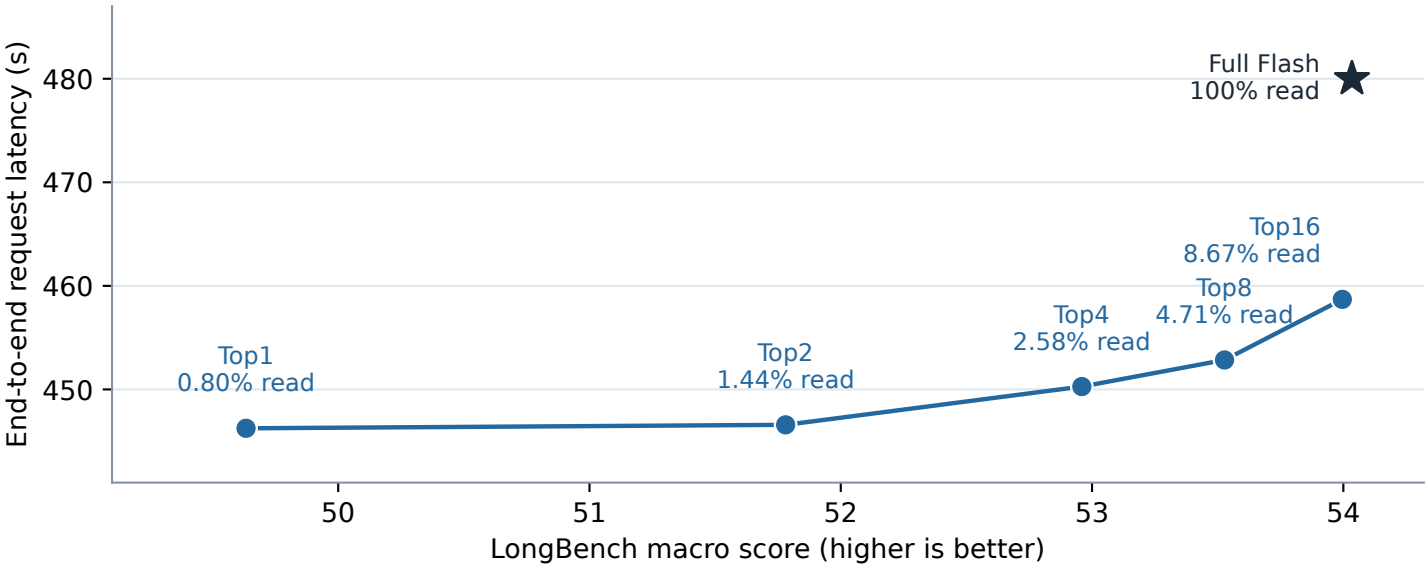

Read labels: measured unique prompt KV reads on routed LongBench decode steps, excluding the generated tail; equal weight across tasks. Full is shown as a reference.
Latency: mean of two 128K requests, including index build and updates, excluding capture.