Shared SAE features across all 16 full-attention layers

27B model · 116 contexts · 64 KV groups · Coverage of at least one shared feature

(a) Top 1% of keys by Full attention vs. remaining 99%

	3	7	11	15	19	23	27	31	35	39	43	47	51	55	59	63
No Q–K · Top 1%	32.7	29.6	35.6	35.7	34.9	35.5	31.9	31.6	32.3	31.2	29.3	35.5	29.6	28.2	25.5	26.4
No Q–K · Rest 99%	5.4	5.5	6.1	5.8	5.8	6.9	7.4	6.5	5.9	5.7	6.0	6.1	5.7	6.6	6.2	5.7
With Q–K · Top 1%	48.5	38.0	40.9	41.6	46.2	42.1	47.9	47.3	44.2	42.6	41.7	47.3	47.2	58.9	62.9	52.5
With Q–K · Rest 99%	4.6	4.5	4.8	4.3	6.1	4.8	6.6	6.0	5.6	5.6	5.6	5.8	5.6	6.0	5.6	4.7

(b) Top 2% of keys by Full attention vs. remaining 98%

	3	7	11	15	19	23	27	31	35	39	43	47	51	55	59	63
No Q–K · Top 2%	27.4	24.9	29.8	29.8	29.3	30.1	27.6	27.4	27.8	27.0	25.6	30.5	25.8	26.0	24.0	23.1
No Q–K · Rest 98%	5.2	5.3	5.9	5.6	5.6	6.7	7.3	6.3	5.7	5.6	5.9	5.9	5.6	6.4	6.0	5.6
With Q–K · Top 2%	39.6	30.8	33.2	33.5	37.8	33.8	39.4	38.9	35.9	35.3	35.6	40.2	39.4	47.3	48.2	41.3
With Q–K · Rest 98%	4.4	4.3	4.5	4.1	5.8	4.5	6.3	5.7	5.4	5.4	5.3	5.5	5.4	5.7	5.3	4.4

Full-attention layer (zero-based)

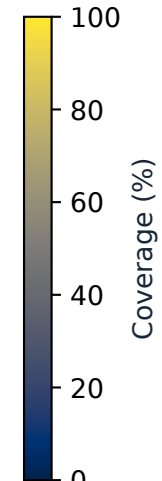

Matched Full Q/K replay for both checkpoints. Each context and KV group receives equal weight.